

Principles of Computer Organization

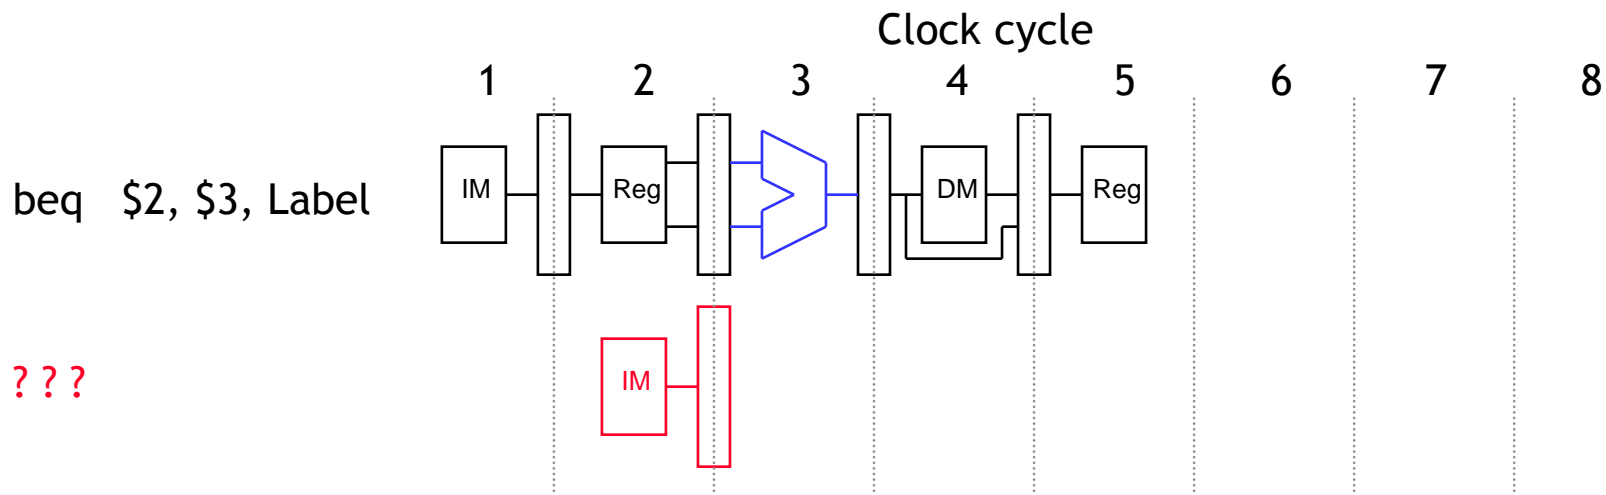
Lecture 12: Pipelining – control hazard and supplementary

Yanyan Liang
Faculty of Information Technology
Macau University of Science and Technology

Control hazard

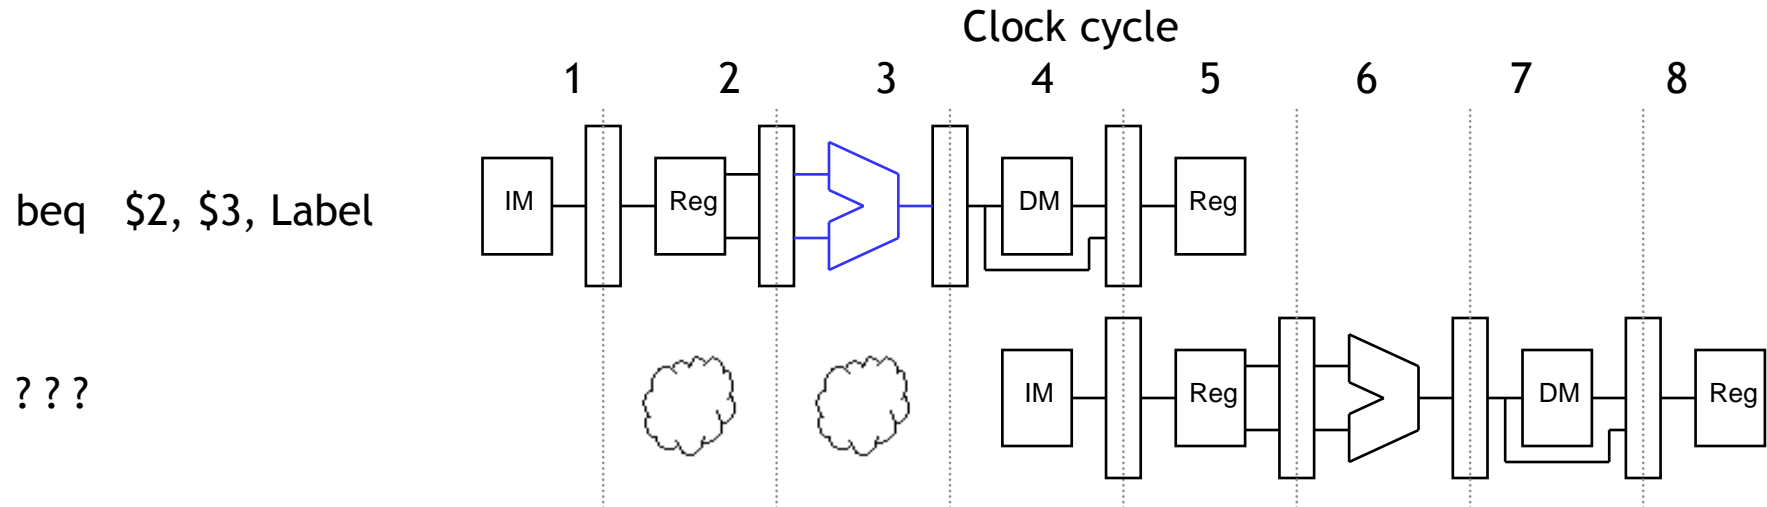
▪ Branch instruction.

- Branch is calculated at EX stage.
- Branch decision is not known until EX is finished, i.e. next instruction cannot be loaded until EX is finished → **control hazard**.



At compile time, insert NOP to stall the pipeline

- Insert two NOPs below BEQ to stall for two cycles.



- But stall often reduces performance of the pipeline.

```
Loop: addi $t1, $t1, 1
```

```
      beq  $t1, $t2, Loop
```


- Assume `$t1 = 0` and `$t2 = 1000` initially. We are expecting the loop can finish within **2000 cycles**, but with branch stalling, it takes **4000 cycles**.

Delayed Branches


- If the branch hardware has been moved to the ID stage, then we can eliminate all branch stalls with **delayed branches** which are defined as always executing the next sequential instruction after the branch instruction – the branch takes effect **after** that next instruction
- MIPS compiler moves an instruction to immediately after the branch that is not affected by the branch (a **safe** instruction) thereby **hiding** the branch delay
- With deeper pipelines, the branch delay grows requiring more than one delay slot
 - 1 Delayed branches have lost popularity compared to more expensive but more flexible (dynamic) hardware branch prediction
 - 1 Growth in available transistors has made hardware branch prediction relatively cheaper

Scheduling Branch Delay Slots

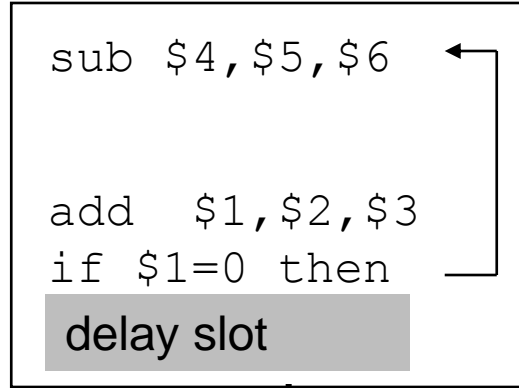
A. From before branch



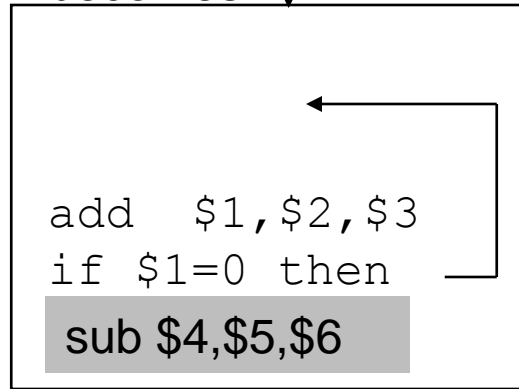
becomes



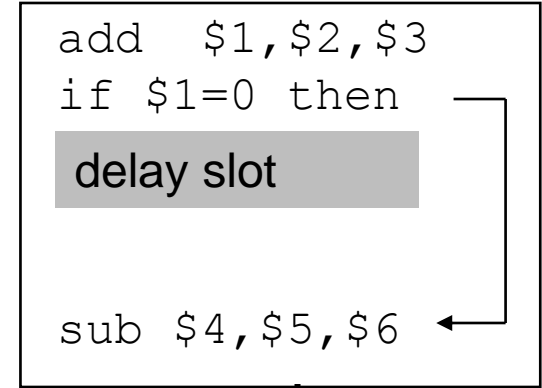
B. From branch target



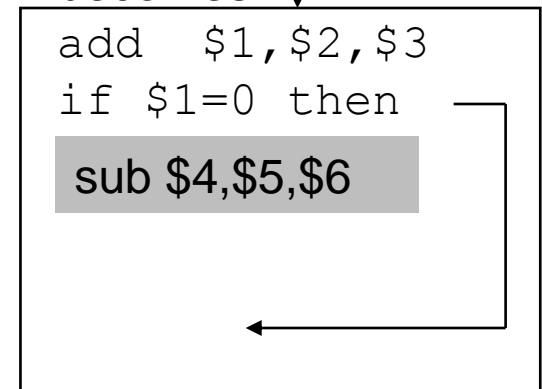
becomes



C. From fall through



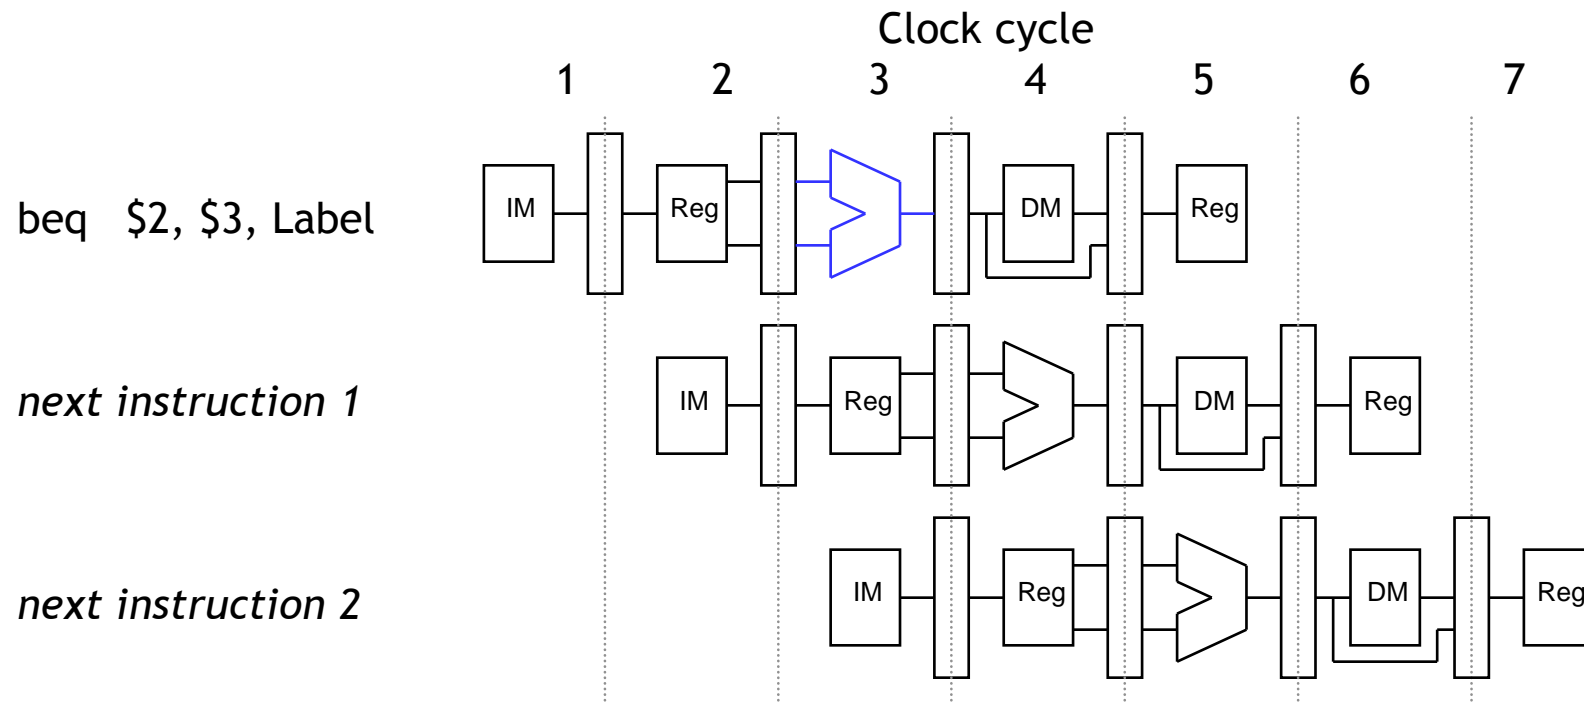
becomes



- A is the best choice, fills delay slot and reduces IC
- In B and C, the `sub` instruction may need to be copied, increasing IC
- In B and C, must be okay to execute `sub` when branch fails

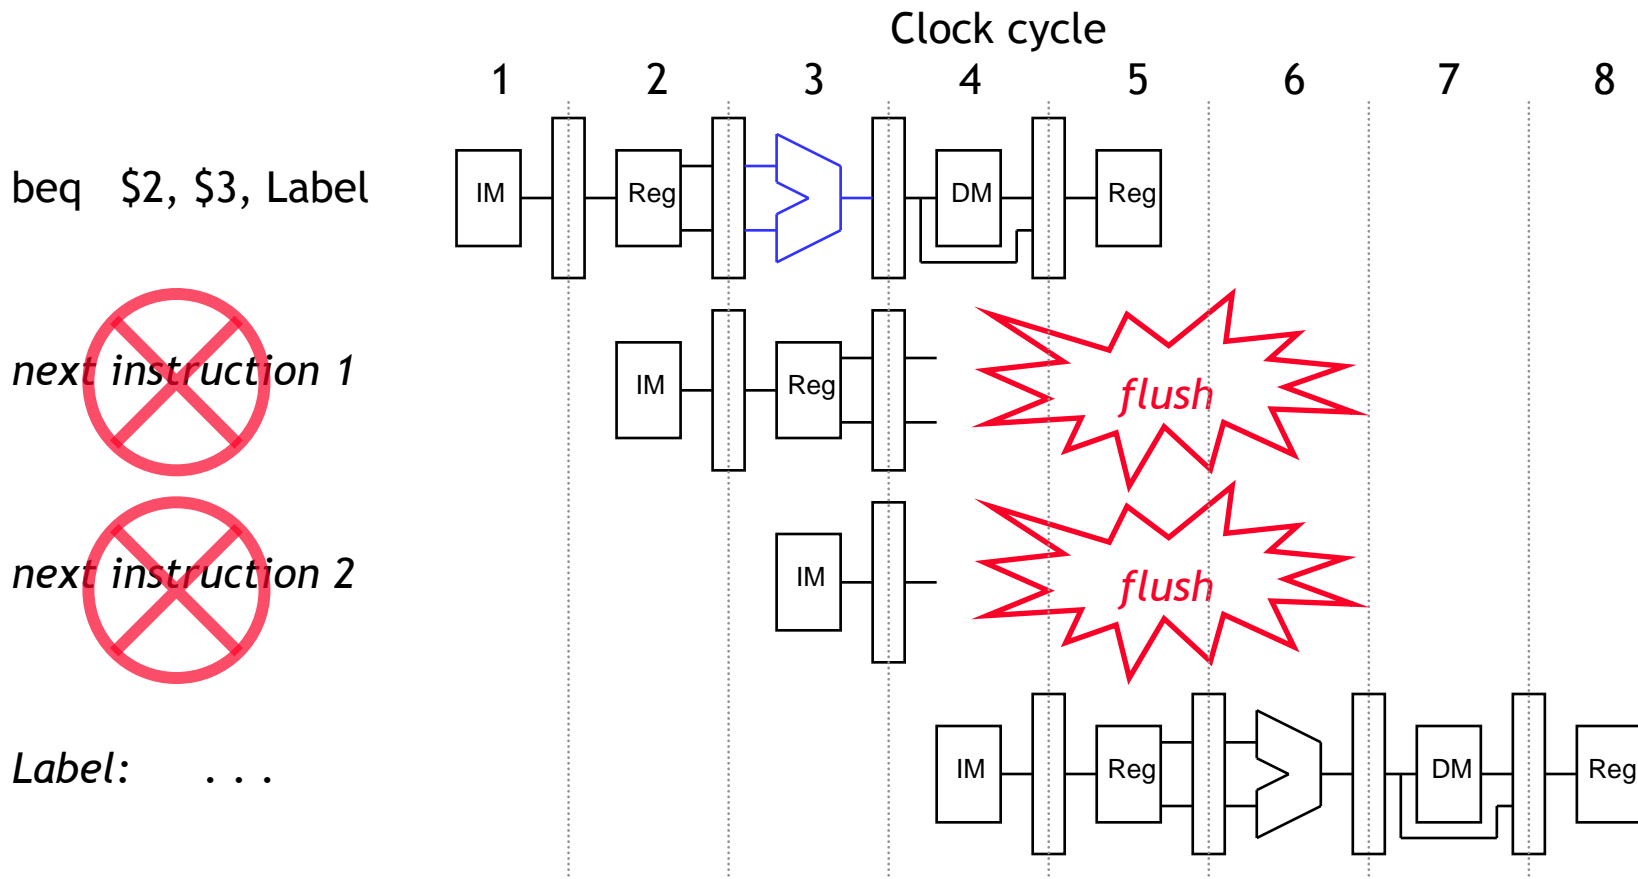
At run time, using branch prediction

- Another approach is to **guess** whether or not the branch is taken.
 - In terms of hardware, it's easier to assume the branch is *not* taken.
 - This way we just increment the PC and continue to execution the next instruction, as for normal instructions.
- If we're correct, then there is no problem and the pipeline keeps going at full speed.



Branch prediction

- If our guess is wrong, that means we have already started executing two incorrect instructions. As a result, we need to discard, or flush, those instructions from the pipeline and begin executing the right one from the branch target address, Label.

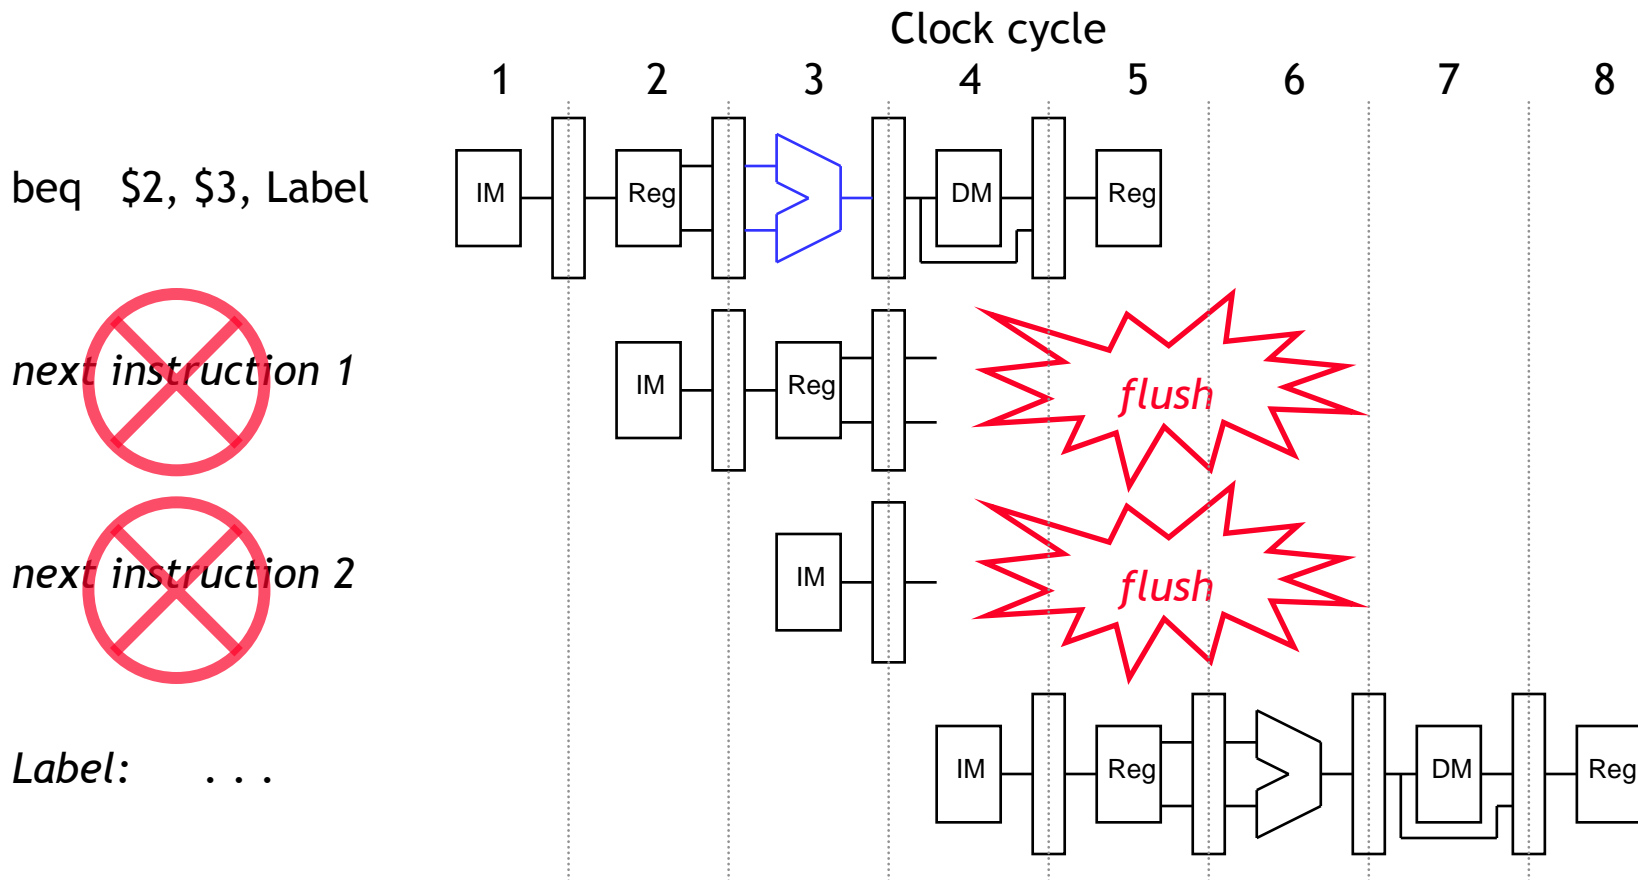


Performance gain of branch prediction

- **In general, branch prediction is worth**
 - **Example: for ($i = 0$; $i < 100$; $i++$)**
if we always predict ($i < 100$) is true \rightarrow 99% correct.
 - **Mispredicting a branch means that two clock cycles are wasted.**
- **A longer pipeline may require more instructions to be flushed for a misprediction, resulting in more wasted time and lower performance.**
 - **If our predictions are just occasionally correct, e.g. 10% accuracy, then it is better to stall and waste two cycles for every branch.**
- **We must be careful that instructions do not modify registers or memory before they get flushed.**

Flush the pipeline: 2 cycles

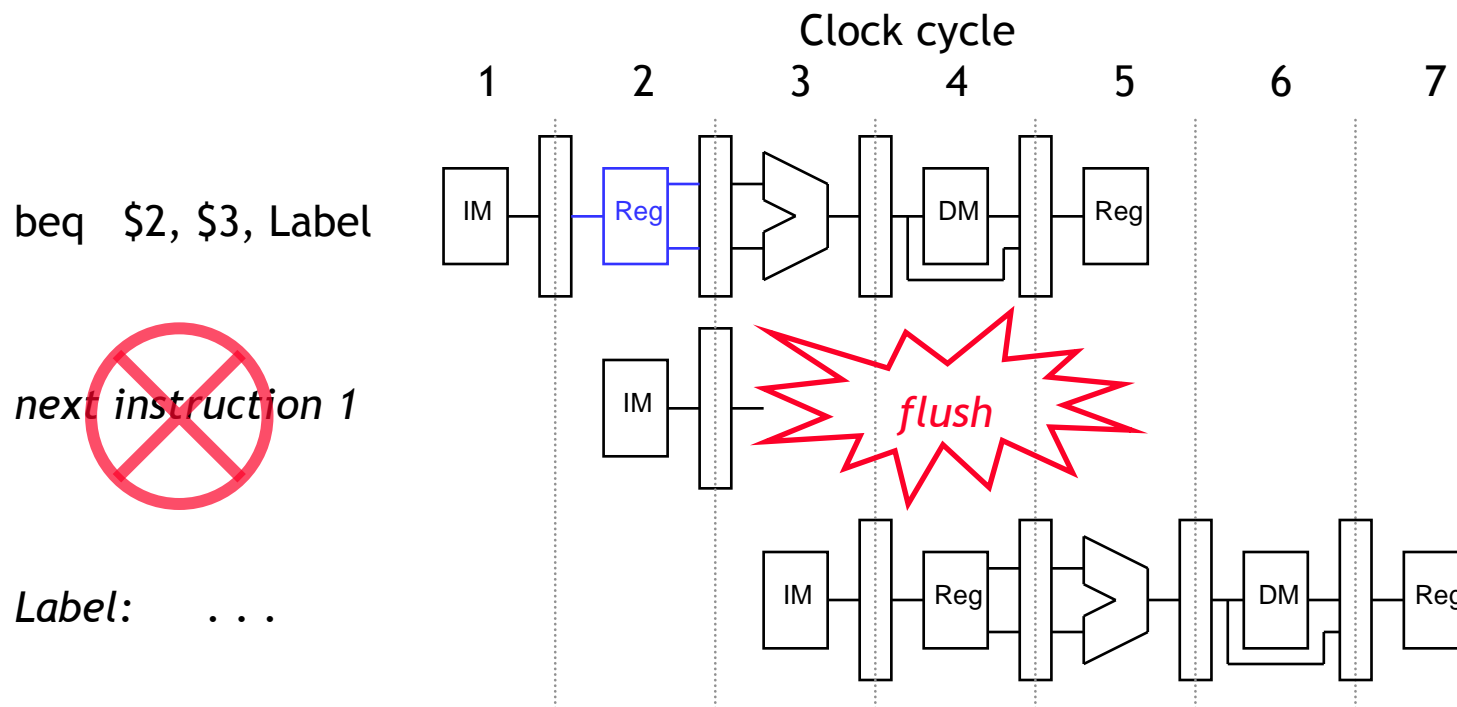
- ALU is used to calculate the branch condition, the condition is only available at **cycle 4**.
- As a result: if misprediction, need to flush **two** instructions.



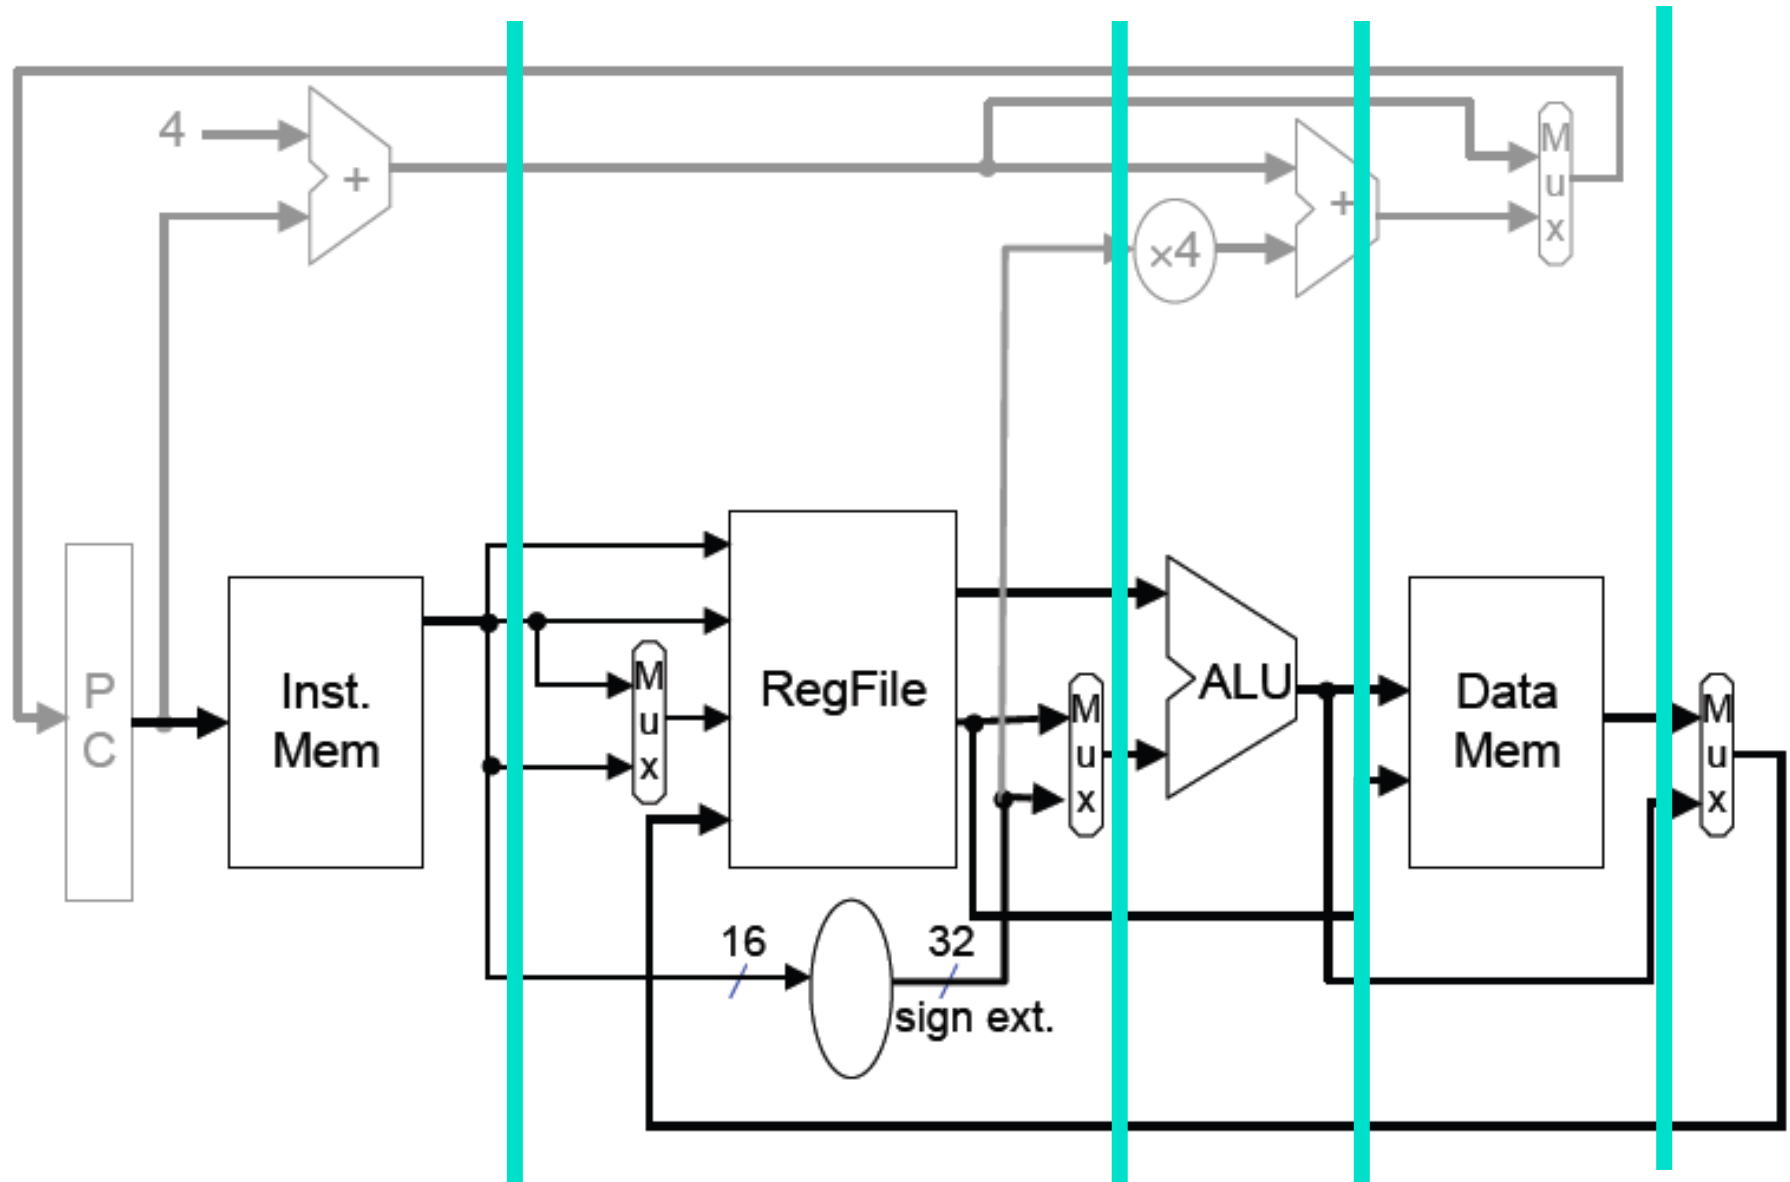
Calculate the branch condition earlier

■ Can we calculate the branch condition earlier?

- Yes. Add extra functional unit to calculate the branch condition at ID stage → The branch condition will be available at **cycle 3**.
- If misprediction, only need to flush **one** instruction.

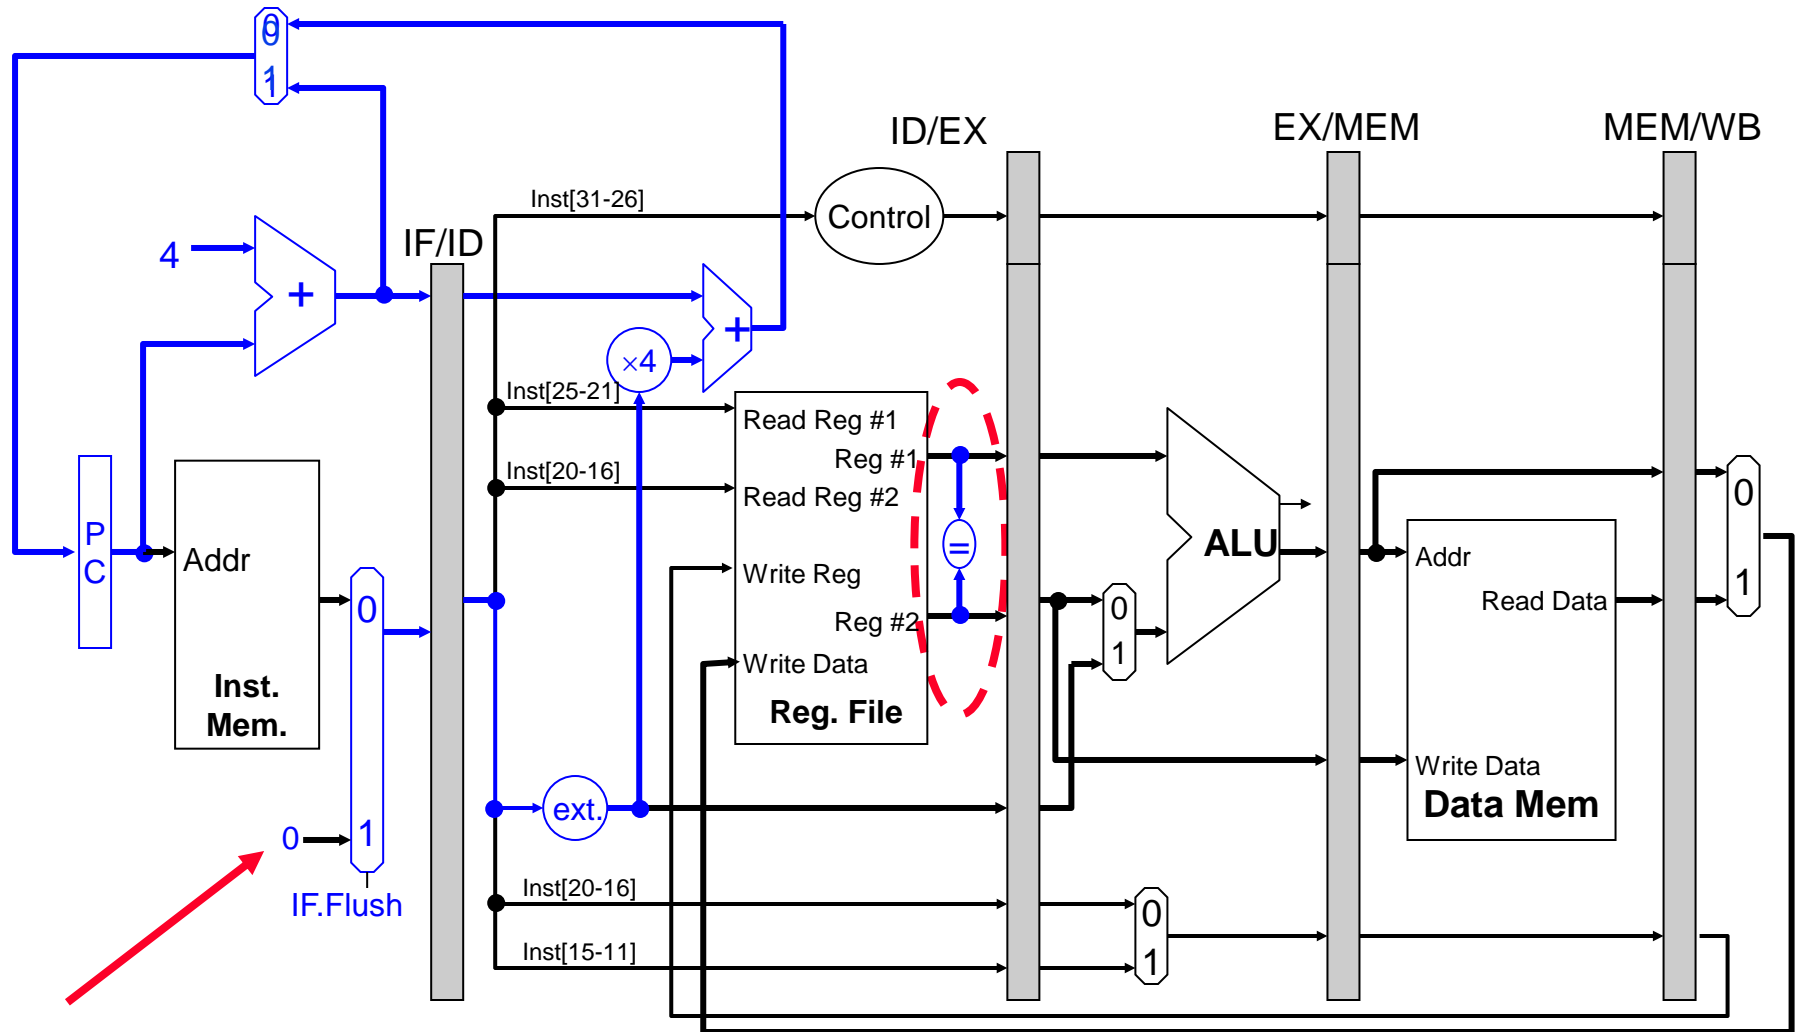


Single cycle processor: partitioning



Branch detection at ID stage

check if two registers are equal.

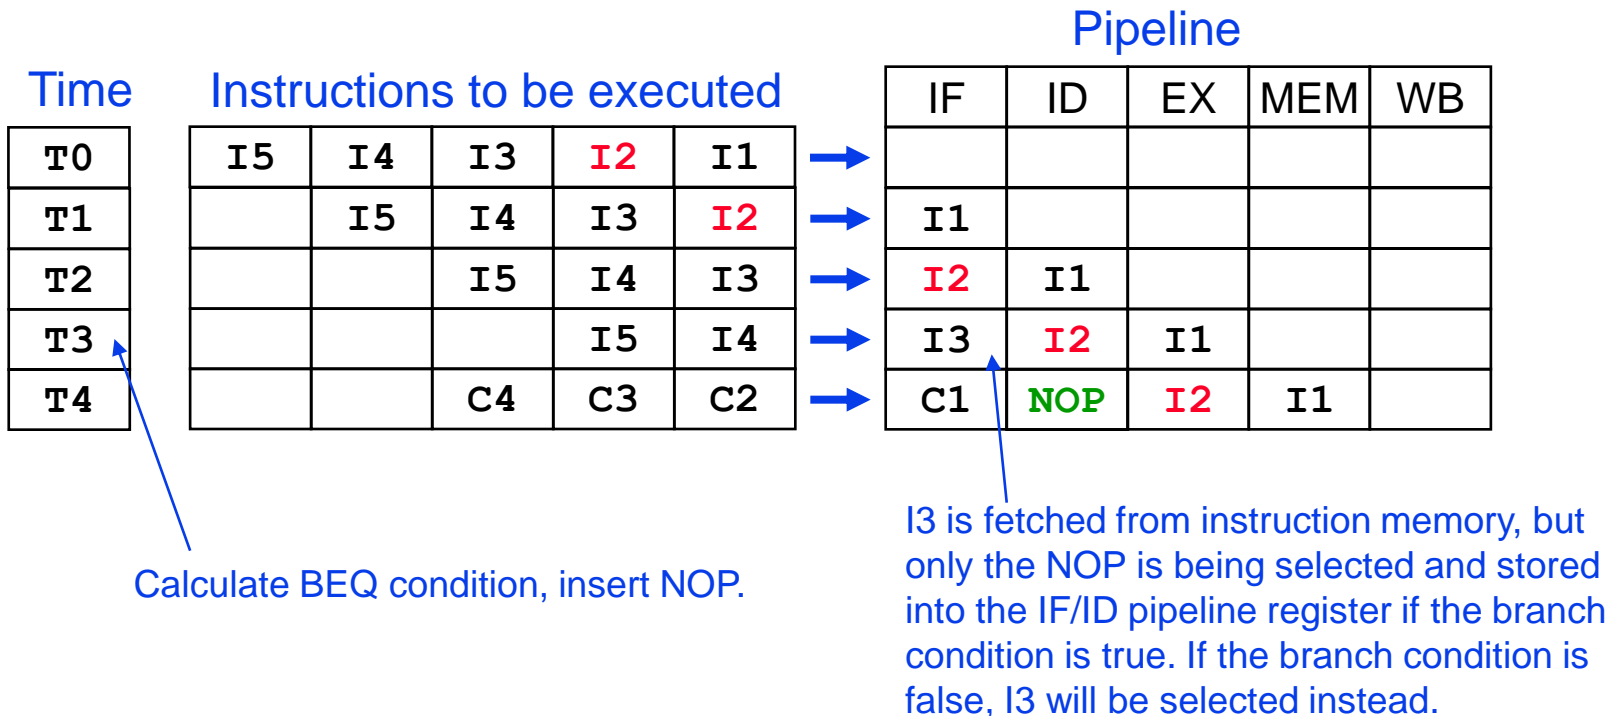


Used to insert NOP instruction into pipeline when there is a branch instruction.

Example: assume I2 is a branch instruction.

▪ How to flush the pipeline when the guess is wrong?

- Assume that our branch prediction strategy is always loading the next instruction (branch not taken). I2 is a beq instruction, C1 is the target instruction.
- Branch condition is calculated at T3, if we find that our prediction is wrong. What should we do?
 - At T3, insert NOP into ID. Fetch C1 at T4.




Example

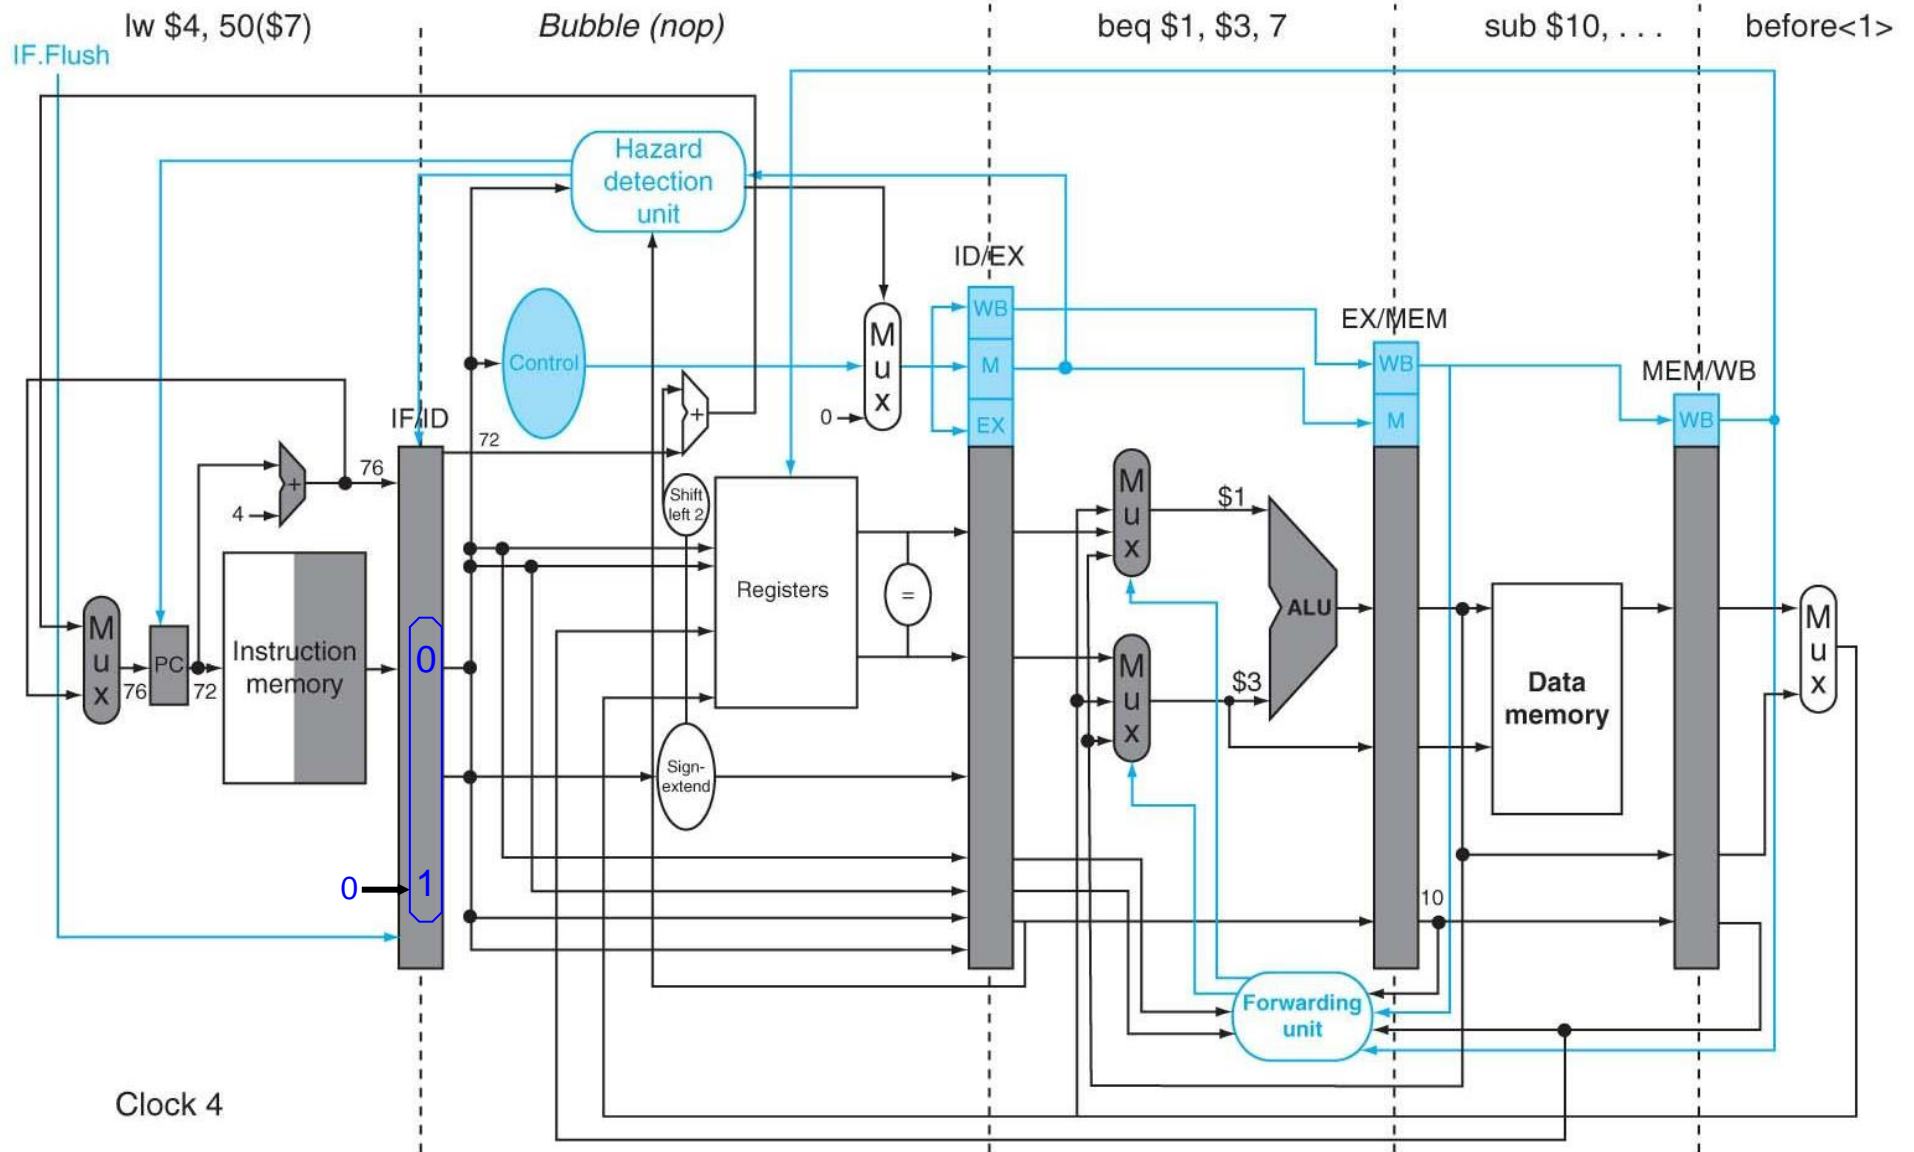
- **Assumption:** The prediction strategy is that branch not taken.

```
36#          sub $10, $4, $8
40#          beq $1, $3, TARGET
44#          and $12, $2, $5
48#          or $13, $2, $6
52#          add $14, $4, $2
56#          slt $15, $6, $7
...
72# TARGET:   lw $4, 50($7)
```

Guess branch not taken: load instruction “and”

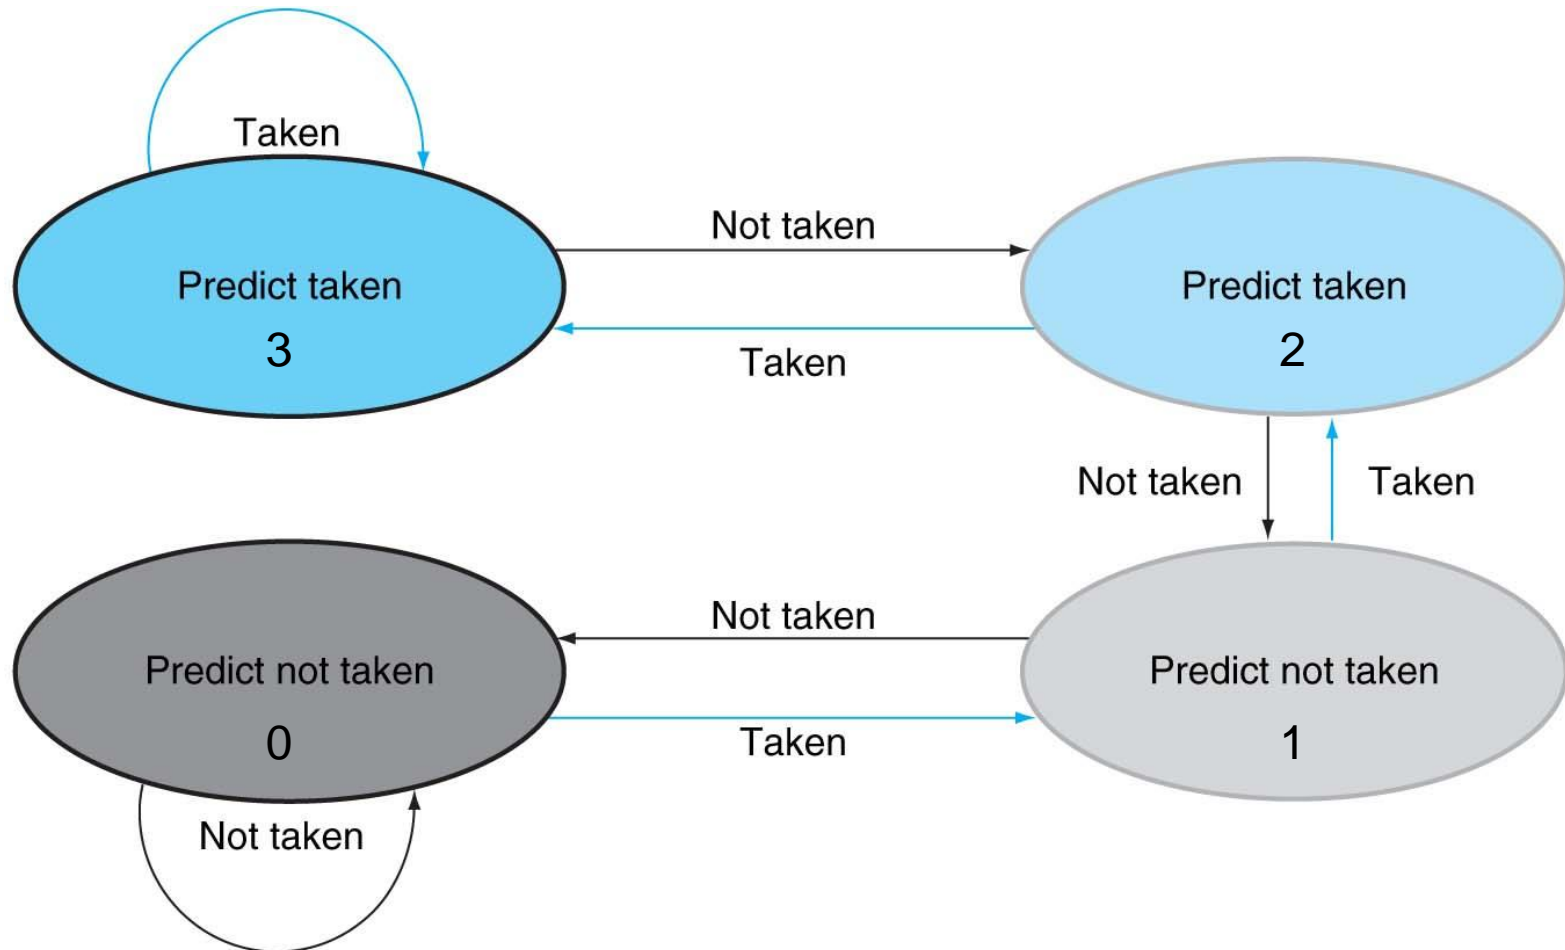


Fetch target instruction “lw”, insert “nop”



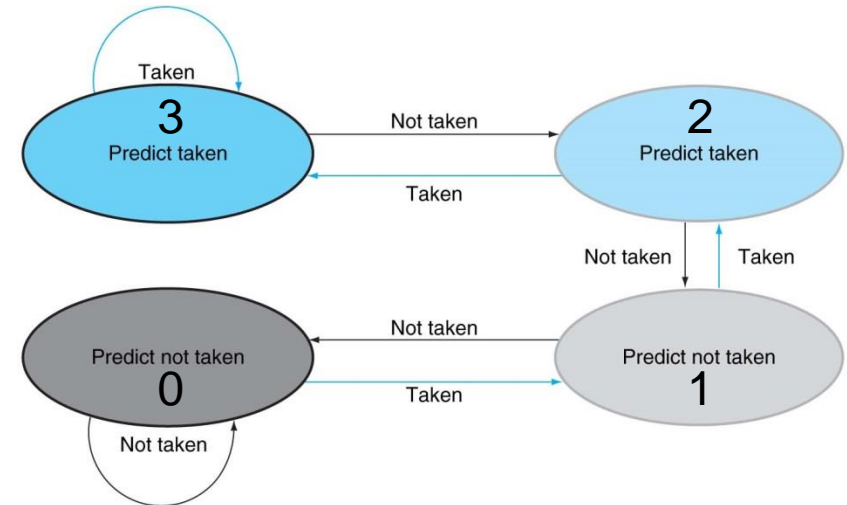
Dynamic prediction (2-bit prediction scheme)

- Predict branch behavior during program execution.



Example of dynamic branch prediction

- Assume that we start off in state 3. A program contains 7 beq instructions.
- Prediction accuracy: $4/7=57\%$



Instruction	Actual execution	Current FSM state	Prediction	Correct?	Next FSM state
beq ...	Taken	3	Taken	Yes	3
beq ...	Not taken	3	Taken	No	2
beq ...	Taken	2	Taken	Yes	3
beq ...	Taken	3	Taken	Yes	3
beq ...	Not Taken	3	Taken	No	2
beq ...	Not Taken	2	Taken	No	1
beq ...	Not Taken	1	Not taken	Yes	0

Summary

- **Control hazard: attempts to make a decision before branch condition is evaluated.**
 - E.g. branch instructions.
 - Solution: insert NOP at compile time; branch prediction with flush.
 - Modify the datapath to detect the branch one cycle earlier.

Exceptions and Interrupts

- **“Unexpected” events requiring change in flow of control**
 - Different ISAs use the terms differently
- **Exception**
 - Arises within the CPU
 - e.g., undefined opcode, overflow, syscall, ...
- **Interrupt**
 - From an external I/O controller
- **Dealing with them without sacrificing performance is hard**

Handling Exceptions

- **In MIPS, exceptions managed by a System Control Coprocessor (CP0)**
- **Save PC of offending (or interrupted) instruction**
 - **In MIPS: Exception Program Counter (EPC)**
- **Save indication of the problem**
 - **In MIPS: Cause register**
 - **We'll assume 1-bit**
 - 0 for undefined opcode, 1 for overflow
- **Jump to handler at 8000 00180**

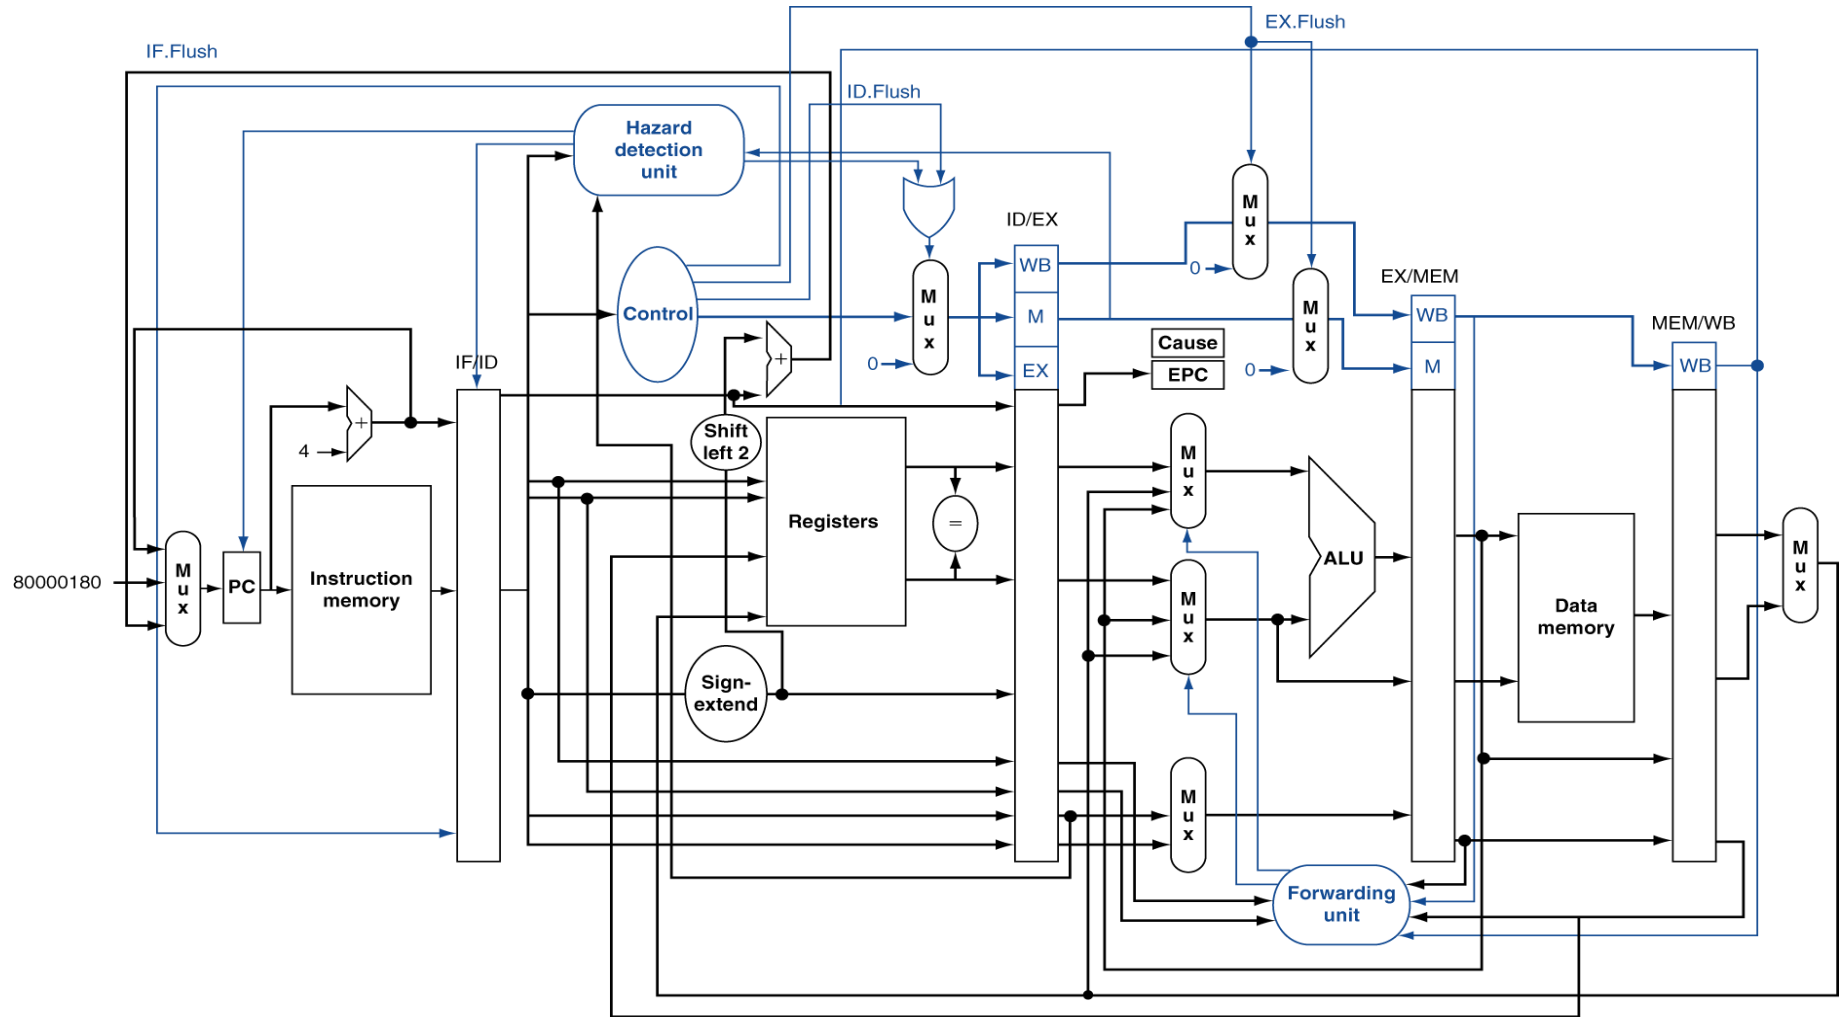
Handler Actions

- **Read cause, and transfer to relevant handler**
- **Determine action required**
- **If restartable**
 - Take corrective action
 - use EPC to return to program
- **Otherwise**
 - Terminate program
 - Report error using EPC, cause, ...

Exceptions in a Pipeline

- **Another form of control hazard**
- **Consider overflow on add in EX stage**
 - add \$1, \$2, \$1
 - Prevent \$1 from being clobbered
 - Complete previous instructions
 - Flush add and subsequent instructions
 - Set Cause and EPC register values
 - Transfer control to handler
- **Similar to mispredicted branch**
 - Use much of the same hardware

Pipeline with Exceptions



Exception Properties

▪ Restartable exceptions

- Pipeline can flush the instruction
- Handler executes, then returns to the instruction
 - Refetched and executed from scratch

▪ PC saved in EPC register

- Identifies causing instruction
- Actually PC + 4 is saved
 - Handler must adjust

Exception Example

▪ Exception on **add** in

40	sub	\$11,	\$2,	\$4
44	and	\$12,	\$2,	\$5
48	or	\$13,	\$2,	\$6
4C	add	\$1,	\$2,	\$1
50	s1t	\$15,	\$6,	\$7
54	lw	\$16,	50(\$7)	

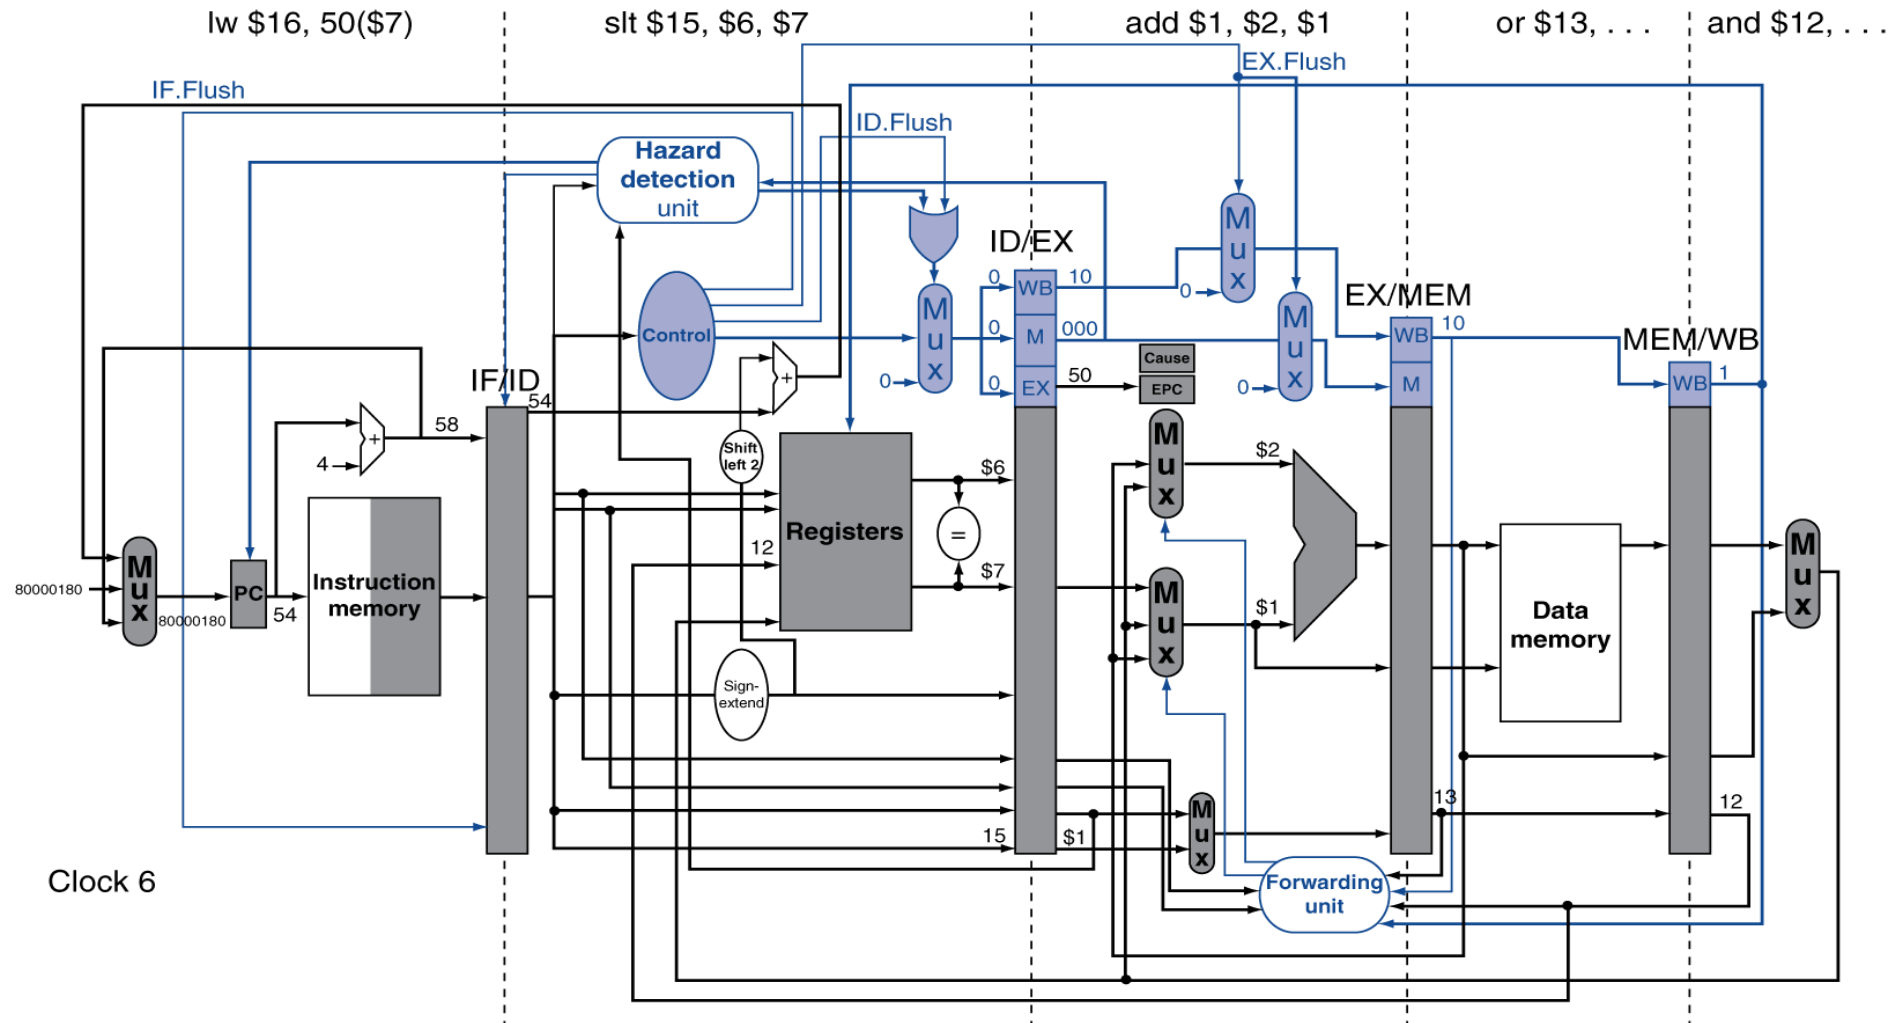
...

▪ Handler

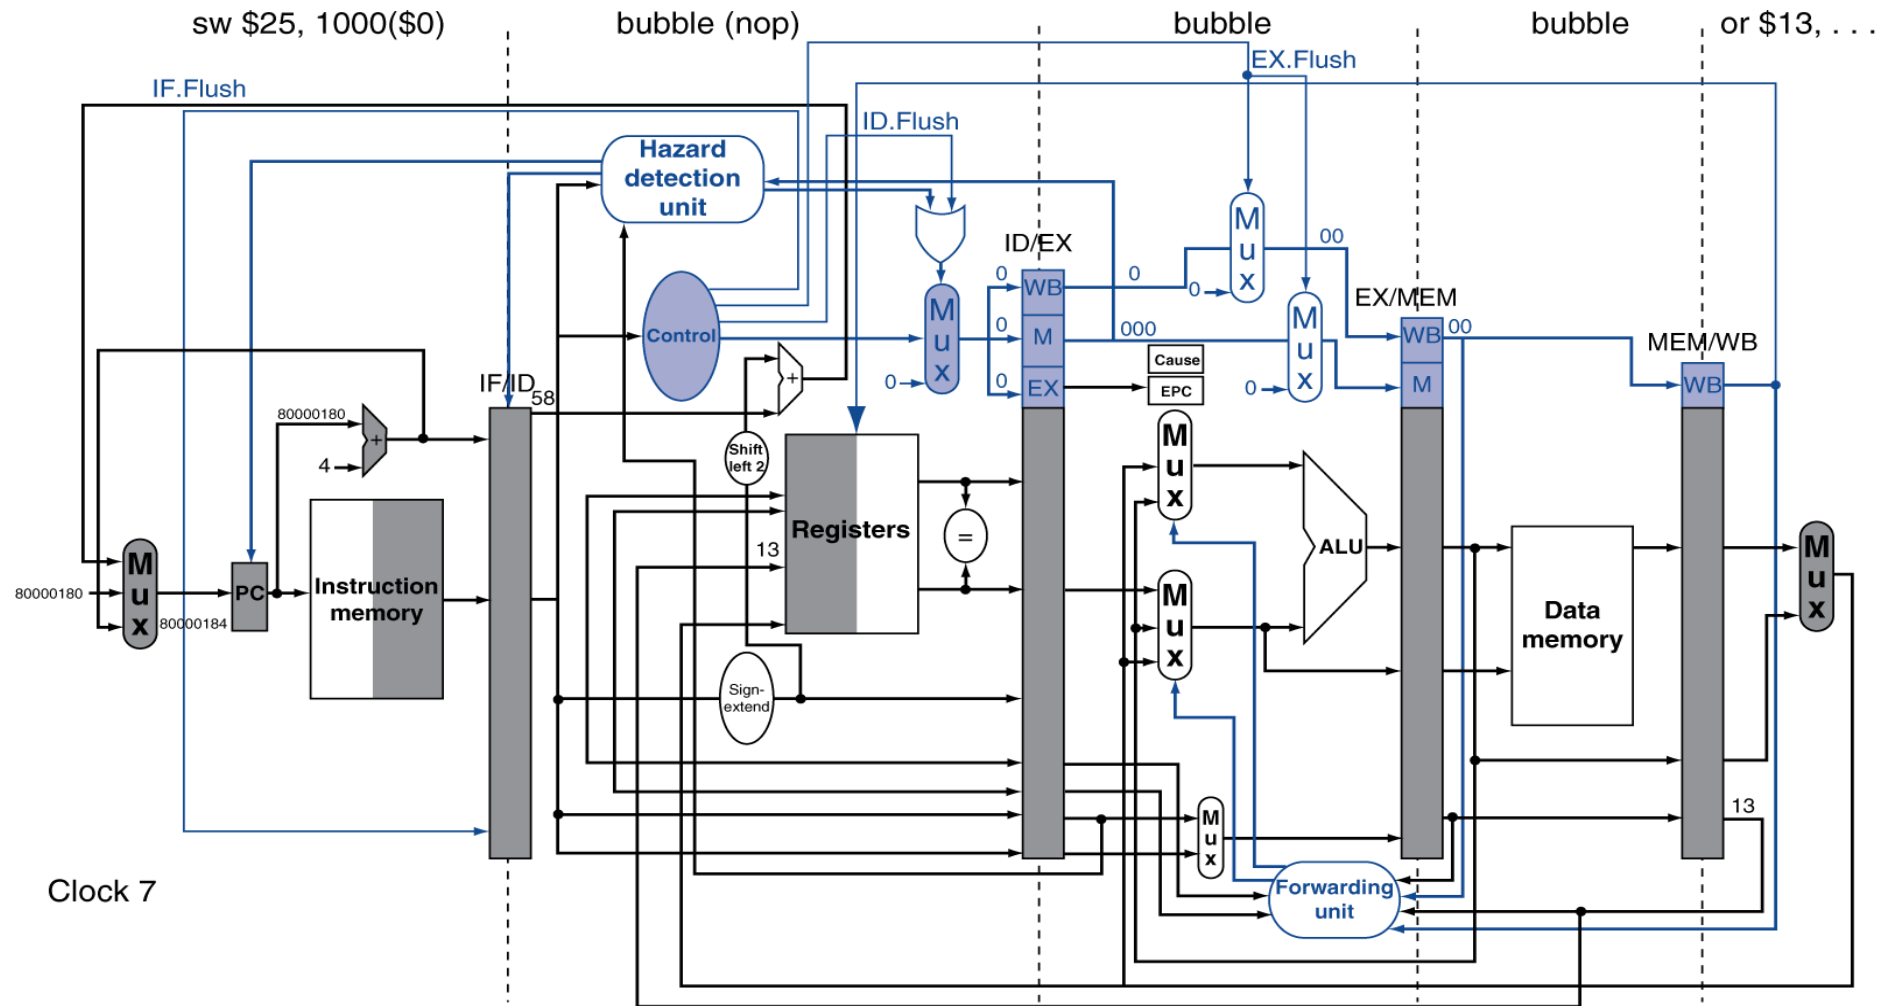
80000180	sw	\$25,	1000(\$0)
80000184	sw	\$26,	1004(\$0)

...

Exception Example



Exception Example



Instruction-Level Parallelism (ILP)

- **Pipelining: executing multiple instructions in parallel**
- **To increase ILP**
 - **Deeper pipeline**
 - Less work per stage \Rightarrow shorter clock cycle
 - **Multiple issue**
 - Replicate pipeline stages \Rightarrow multiple pipelines
 - Start multiple instructions per clock cycle
 - $CPI < 1$, so use Instructions Per Cycle (IPC)
 - E.g., 4GHz 4-way multiple-issue
 - 16 BIPS, peak $CPI = 0.25$, peak $IPC = 4$
 - But dependencies reduce this in practice

Multiple Issue

▪ Static multiple issue

- Compiler groups instructions to be issued together
- Packages them into “issue slots”
- Compiler detects and avoids hazards

▪ Dynamic multiple issue

- CPU examines instruction stream and chooses instructions to issue each cycle
- Compiler can help by reordering instructions
- CPU resolves hazards using advanced techniques at runtime

Speculation

- **“Guess” what to do with an instruction**
 - Start operation as soon as possible
 - Check whether guess was right
 - If so, complete the operation
 - If not, roll-back and do the right thing
- **Common to static and dynamic multiple issue**
- **Examples**
 - Speculate on branch outcome
 - Roll back if path taken is different
 - Speculate on load
 - Roll back if location is updated

Compiler/Hardware Speculation

- **Compiler can reorder instructions**
 - e.g., move load before branch
 - Can include “fix-up” instructions to recover from incorrect guess
- **Hardware can look ahead for instructions to execute**
 - Buffer results until it determines they are actually needed
 - Flush buffers on incorrect speculation

Speculation and Exceptions

- **What if exception occurs on a speculatively executed instruction?**
 - e.g., speculative load before null-pointer check
- **Static speculation**
 - Can add ISA support for deferring exceptions
- **Dynamic speculation**
 - Can buffer exceptions until instruction completion (which may not occur)

Static Multiple Issue

- **Compiler groups instructions into “issue packets”**
 - Group of instructions that can be issued on a single cycle
 - Determined by pipeline resources required
- **Think of an issue packet as a very long instruction**
 - Specifies multiple concurrent operations
 - \Rightarrow Very Long Instruction Word (VLIW)

Scheduling Static Multiple Issue

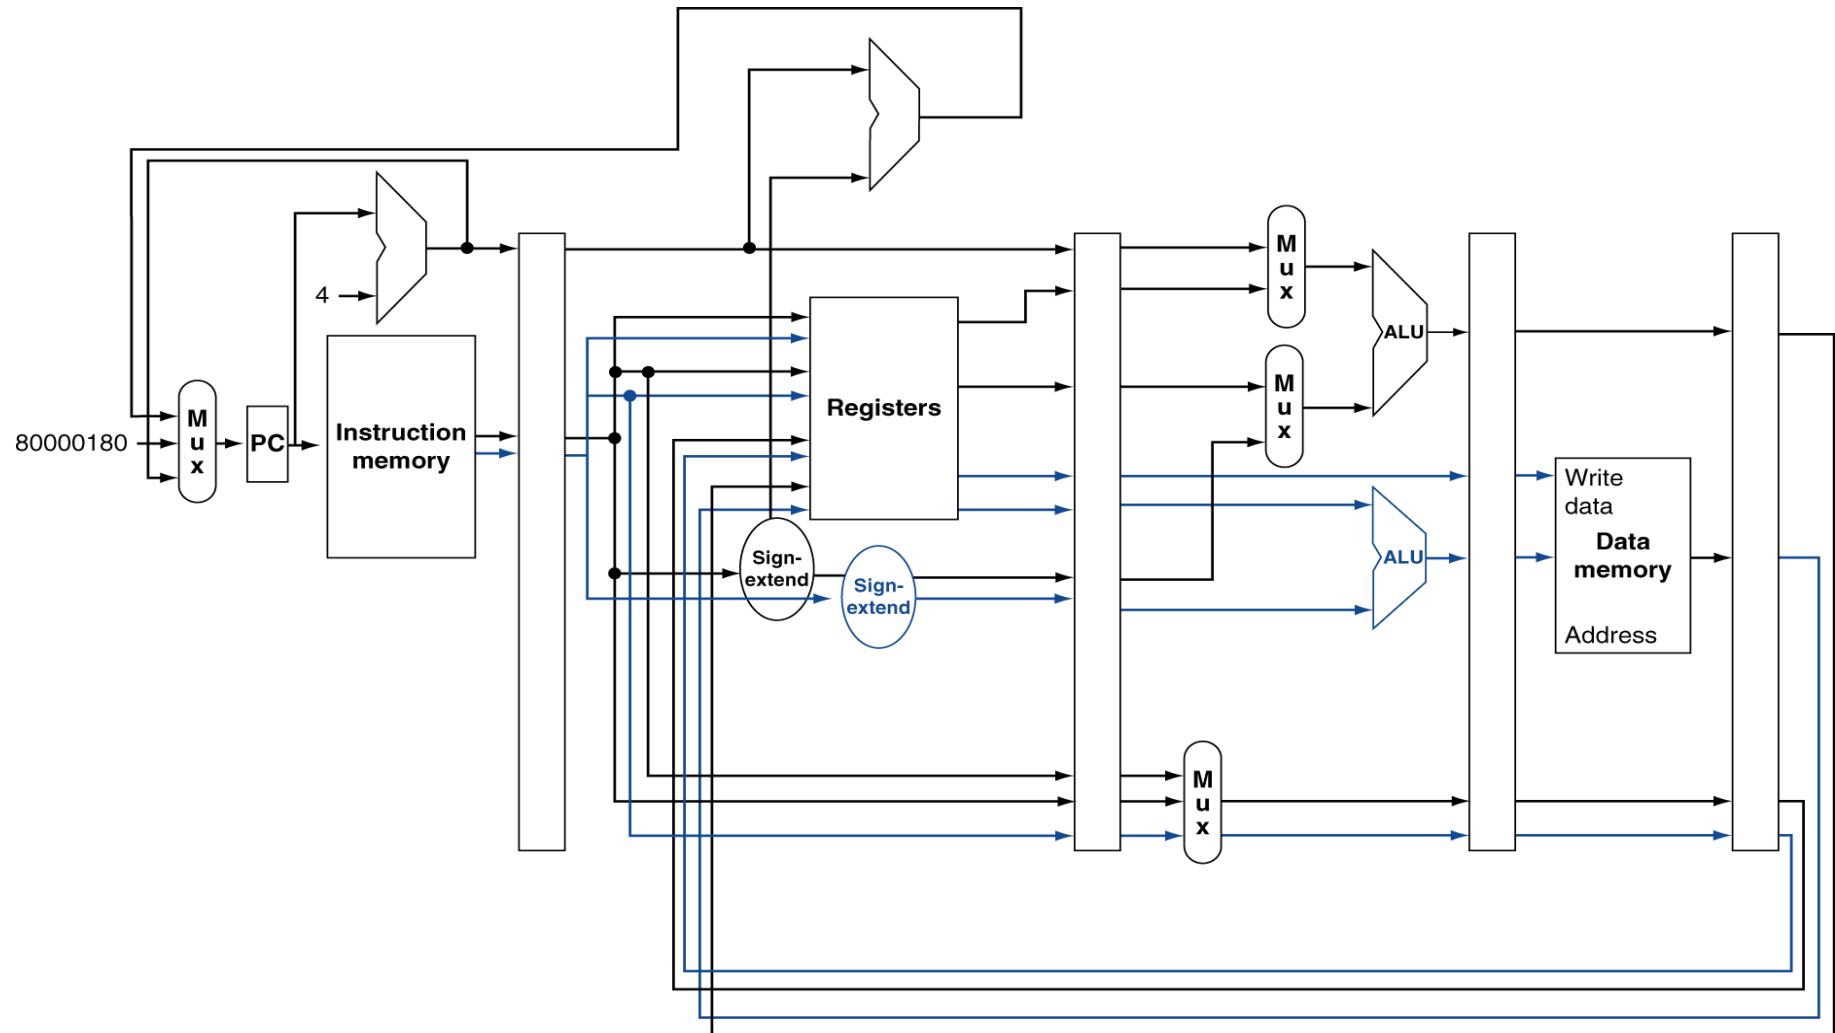
- **Compiler must remove some/all hazards**
 - Reorder instructions into issue packets
 - No dependencies with a packet
 - Possibly some dependencies between packets
 - Varies between ISAs; compiler must know!
 - Pad with nop if necessary

MIPS with Static Dual Issue

- **Two-issue packets**
 - One ALU/branch instruction
 - One load/store instruction
 - 64-bit aligned
 - ALU/branch, then load/store
 - Pad an unused instruction with nop

Address	Instruction type	Pipeline Stages						
n	ALU/branch	IF	ID	EX	MEM	WB		
n + 4	Load/store	IF	ID	EX	MEM	WB		
n + 8	ALU/branch		IF	ID	EX	MEM	WB	
n + 12	Load/store		IF	ID	EX	MEM	WB	
n + 16	ALU/branch			IF	ID	EX	MEM	WB
n + 20	Load/store			IF	ID	EX	MEM	WB

MIPS with Static Dual Issue



Hazards in the Dual-Issue MIPS

- More instructions executing in parallel
- EX data hazard
 - Forwarding avoided stalls with single-issue
 - Now can't use ALU result in load/store in same packet
 - add `$t0`, `$s0`, `$s1`
load `$s2`, 0(`$t0`)
 - Split into two packets, effectively a stall
- Load-use hazard
 - Still one cycle use latency, but now two instructions
- More aggressive scheduling required

Scheduling Example

▪ Schedule this for dual-issue MIPS

```
Loop: lw    $t0, 0($s1)      # $t0=array element
      addu  $t0, $t0, $s2    # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi  $s1, $s1, -4     # decrement pointer
      bne   $s1, $zero, Loop # branch $s1!=0
```

	ALU/branch	Load/store	cycle
Loop:	nop	lw \$t0, 0(\$s1)	1
	addi \$s1, \$s1, -4	nop	2
	addu \$t0, \$t0, \$s2	nop	3
	bne \$s1, \$zero, Loop	sw \$t0, 4(\$s1)	4

■ $IPC = 5/4 = 1.25$ (c.f. peak $IPC = 2$)

Loop Unrolling

- **Replicate loop body to expose more parallelism**
 - Reduces loop-control overhead
- **Use different registers per replication**
 - Called “register renaming”
 - Avoid loop-carried “anti-dependencies”
 - Store followed by a load of the same register
 - Aka “name dependence”
 - Reuse of a register name

Loop Unrolling Example

	ALU/branch	Load/store	cycle
Loop:	addi \$s1 , \$s1, -16	lw \$t0 , 0(\$s1)	1
	nop	lw \$t1 , 12(\$s1)	2
	addu \$t0 , \$t0 , \$s2	lw \$t2 , 8(\$s1)	3
	addu \$t1 , \$t1 , \$s2	lw \$t3 , 4(\$s1)	4
	addu \$t2 , \$t2 , \$s2	sw \$t0 , 16(\$s1)	5
	addu \$t3 , \$t4 , \$s2	sw \$t1 , 12(\$s1)	6
	nop	sw \$t2 , 8(\$s1)	7
	bne \$s1 , \$zero, Loop	sw \$t3 , 4(\$s1)	8

▪ **IPC = 14/8 = 1.75**

- **Closer to 2, but at cost of registers and code size**

Dynamic Multiple Issue

- **“Superscalar” processors**
- **CPU decides whether to issue 0, 1, 2, ... each cycle**
 - Avoiding structural and data hazards
- **Avoids the need for compiler scheduling**
 - Though it may still help
 - Code semantics ensured by the CPU

Dynamic Pipeline Scheduling

- **Allow the CPU to execute instructions out of order to avoid stalls**

- But commit result to registers in order

- **Example**

```
lw      $t0, 20($s2)
addu    $t1, $t0, $t2
sub      $s4, $s4, $t3
slli    $t5, $s4, 20
```

- Can start sub while addu is waiting for lw

Why Do Dynamic Scheduling?

- **Why not just let the compiler schedule code?**
- **Not all stalls are predicable**
 - e.g., cache misses
- **Can't always schedule around branches**
 - Branch outcome is dynamically determined
- **Different implementations of an ISA have different latencies and hazards**

Does Multiple Issue Work?

The BIG Picture

- **Yes, but not as much as we'd like**
- **Programs have real dependencies that limit ILP**
- **Some dependencies are hard to eliminate**
 - e.g., pointer aliasing
- **Some parallelism is hard to expose**
 - Limited window size during instruction issue
- **Memory delays and limited bandwidth**
 - Hard to keep pipelines full
- **Speculation can help if done well**

```
int i;  
int *pi = &i;
```

pi alias i

```
int i;  
void foo(int &i1, int &i2){}  
foo(i,i);
```

i1 alias i2

Power Efficiency

- **Complexity of dynamic scheduling and speculations requires power**
- **Multiple simpler cores may be better**

Microprocessor	Year	Clock Rate	Pipeline Stages	Issue width	Out-of-order/ Speculation	Cores	Power
i486	1989	25MHz	5	1	No	1	5W
Pentium	1993	66MHz	5	2	No	1	10W
Pentium Pro	1997	200MHz	10	3	Yes	1	29W
P4 Willamette	2001	2000MHz	22	3	Yes	1	75W
P4 Prescott	2004	3600MHz	31	3	Yes	1	103W
Core	2006	2930MHz	14	4	Yes	2	75W
UltraSparc III	2003	1950MHz	14	4	No	1	90W
UltraSparc T1	2005	1200MHz	6	1	No	8	70W

Concluding Remarks

- **ISA influences design of datapath and control**
- **Datapath and control influence design of ISA**
- **Pipelining improves instruction throughput using parallelism**
 - **More instructions completed per second**
 - **Latency for each instruction not reduced**
- **Hazards: structural, data, control**
- **Multiple issue and dynamic scheduling (ILP)**
 - **Dependencies limit achievable parallelism**
 - **Complexity leads to the power wall**